

Ingestion Recovery Manual

Page 1 - Failure detection

The ingestion pipeline parses and chunks a complete document before requesting embeddings. Chroma is updated only after every vector in the document batch has been validated.

If an embedding request fails, the pipeline reports the upstream error and performs no partial write. Operators should record the document ID, failed batch index, and source filename before retrying.

```
failure_state = EMBEDDING_FAILED
```

Recovery continues on the next page with a stable retry token. The document must not be edited between failure detection and retry, because a content change produces a different document ID.

Ingestion Recovery Manual

Page 2 - Retry token

Create the recovery marker `INGEST_RETRY_TOKEN` from the original document ID and failed batch index. Reuse this token for every retry of the same failed ingestion attempt.

```
INGEST_RETRY_TOKEN = document_id + ':' + batch_index
```

A retry must use the same embedding model, vector dimension, chunk strategy, chunk size, and overlap. Changing these parameters starts a new ingestion representation rather than recovering the failed one.

After all vectors are returned, the pipeline may upsert the complete chunk list. Verification continues on page 3 and must be completed before the recovery marker is cleared.

Ingestion Recovery Manual

Page 3 - Verification

Compare the stored chunk count with the chunk count returned by the ingestion result. The values must match for the recovered document.

Inspect at least one stored chunk from every PDF page. Each chunk must retain its source filename, source page metadata, content hash, and the chunk strategy used during recovery.

Finally, run a dense retrieval query for INGEST_RETRY_TOKEN. Recovery is complete only when the relevant page 2 chunk is returned and the stored chunk count plus source page metadata checks both pass.

```
verification_status = RECOVERY_COMPLETE
```